

# VASP on GPUs

The background of the slide features a large, faint, dark red crest of the University of Chicago. The crest is a shield-shaped emblem. At the top, it contains an open book with the Latin motto 'Crescat Scientia Vita Excolatur' (Let knowledge grow, life be cultivated). Below the book is a bald eagle with its wings spread, perched on a branch. The entire crest is rendered in a dark red color that blends with the slide's background.

Max Hutchinson  
University of Chicago

May 11, 2016



# Big thanks to

## Carnegie Mellon group

- Michael Widom

## ENS/IFPEN group

- Paul Fleurat-Lessard
- Thomas Guignon
- Ani Anciaux-Sedrakian
- Philippe Sautet

## RWTH Aachen Group

- Stefan Maintz
- Bernhard Eck
- Richard Dronskowski



# Big thanks to

## University of Vienna group

■ Georg Kresse

■ Martijn Marsman

■ Doris Vogtenhuber

## NVIDIA

■ Christoph Angerer

■ Sarah Tariq

■ Anthony Scudiero

■ Jeroen Bédorf

■ Dusan Stosic

■ Darko Stosic

■ Arash Ashari

■ Paul Springer

■ Przemek Tredak

■ Mark Berger

■ Jerry Chen

■ Cliff Woolley



# A brief history

## Multiple prototypes (2009-2012)

- Diagonalization for traditional DFT<sup>12</sup>(IFPEN, ENS, Aachen)
- Exact-exchange for hybrid functionals<sup>3</sup>(CMU, UChicago)

## Cooperation and tuning (2012 - 2014)

- Merge prototypes with VASP 5.3.1
- Performance tune with NVIDIA engineers

---

<sup>1</sup>M. Hacene et al., DOI:10.1002/jcc.23096

<sup>2</sup>S. Maintz et al., DOI:10.1016/j.cpc.2011.03.010

<sup>3</sup>M. Hutchinson and M. Widom, DOI:10.1016/j.cpc.2012.02.017



# A brief history

## Multiple prototypes (2009-2012)

- Diagonalization for traditional DFT<sup>12</sup>(IFPEN, ENS, Aachen)
- Exact-exchange for hybrid functionals<sup>3</sup>(CMU, UChicago)

## Cooperation and tuning (2012 - 2014)

- Merge prototypes with VASP 5.3.1
- Performance tune with NVIDIA engineers

---

<sup>1</sup>M. Hacene et al., DOI:10.1002/jcc.23096

<sup>2</sup>S. Maintz et al., DOI:10.1016/j.cpc.2011.03.010

<sup>3</sup>M. Hutchinson and M. Widom, DOI:10.1016/j.cpc.2012.02.017



# A brief history

## Multiple prototypes (2009-2012)

- Diagonalization for traditional DFT<sup>12</sup>(IFPEN, ENS, Aachen)
- Exact-exchange for hybrid functionals<sup>3</sup>(CMU, UChicago)

## Cooperation and tuning (2012 - 2014)

- Merge prototypes with VASP 5.3.1
- Performance tune with NVIDIA engineers

<sup>1</sup>M. Hacene et al., DOI:10.1002/jcc.23096

<sup>2</sup>S. Maintz et al., DOI:10.1016/j.cpc.2011.03.010

<sup>3</sup>M. Hutchinson and M. Widom, DOI:10.1016/j.cpc.2012.02.017

# A brief history



## Acceptance and distribution (2015)

- GPU support accepted by Vienna
- Integrated development environments
- **Established correctness**
- **Included in standard VASP release**



# Establishing correctness

We've taken a three-pronged approach to validation:

1. Internal testing against  $\sim 50$  cases collected from collaborators
  - Focus on actively ported algorithms and models
2. Acceptance testing against  $\sim 100$  cases by Vienna
  - Cover wider variety of VASP usage patterns
3. Beta testing by 37 early access groups
  - Cover a wider variety of hardware and environments





# Establishing correctness

We've taken a three-pronged approach to validation:

1. Internal testing against  $\sim 50$  cases collected from collaborators
  - Focus on actively ported algorithms and models
2. Acceptance testing against  $\sim 100$  cases by Vienna
  - Cover wider variety of VASP usage patterns
3. Beta testing by 37 early access groups
  - Cover a wider variety of hardware and environments



# Establishing correctness

We've taken a three-pronged approach to validation:

1. Internal testing against  $\sim 50$  cases collected from collaborators
  - Focus on actively ported algorithms and models
2. Acceptance testing against  $\sim 100$  cases by Vienna
  - Cover wider variety of VASP usage patterns
3. Beta testing by 37 early access groups
  - Cover a wider variety of hardware and environments



# Establishing correctness

We've taken a three-pronged approach to validation:

1. Internal testing against  $\sim 50$  cases collected from collaborators
  - Focus on actively ported algorithms and models
2. Acceptance testing against  $\sim 100$  cases by Vienna
  - Cover wider variety of VASP usage patterns
3. Beta testing by 37 early access groups
  - Cover a wider variety of hardware and environments



# Beta testing

## Three types of issues

- Use of unsupported features
- Merge with site-customized files (esp. main.F)
- Bugs in edge cases

## Generally positive feedback

- “The short version is ‘it works’”
- “So far I found no problems, the code is fast and stable.”
- “Absolute time to solution is faster with GPUs.”



# Feature support

## Fully supported

- Davidson
- RMM-DIIS
- Exact-exchange
- R-space projection
- Non-collinear
- KPAR

## Passively supported

- [sc]GW[0]
- Damped
- All (Algo)

## Unsupported

- G-space projection
- NCORE > 1
- EFIELD\_PEAD



# Feature support

## Fully supported

- |                      |                 |                  |
|----------------------|-----------------|------------------|
| ■ Davidson           | ■ RMM-DIIS      | ■ Exact-exchange |
| ■ R-space projection | ■ Non-collinear | ■ KPAR           |

## Passively supported

- |             |          |              |
|-------------|----------|--------------|
| ■ [sc]GW[0] | ■ Damped | ■ All (Algo) |
|-------------|----------|--------------|

## Unsupported

- |                      |             |               |
|----------------------|-------------|---------------|
| ■ G-space projection | ■ NCORE > 1 | ■ EFIELD_PEAD |
|----------------------|-------------|---------------|



# Feature support

## Fully supported

- |                      |                 |                  |
|----------------------|-----------------|------------------|
| ■ Davidson           | ■ RMM-DIIS      | ■ Exact-exchange |
| ■ R-space projection | ■ Non-collinear | ■ KPAR           |

## Passively supported

- |             |          |              |
|-------------|----------|--------------|
| ■ [sc]GW[0] | ■ Damped | ■ All (Algo) |
|-------------|----------|--------------|

## Unsupported

- |                      |             |               |
|----------------------|-------------|---------------|
| ■ G-space projection | ■ NCORE > 1 | ■ EFIELD_PEAD |
|----------------------|-------------|---------------|



# Feature support

## Fully supported

- Davidson
- RMM-DIIS
- Exact-exchange
- R-space projection
- Non-collinear
- KPAR

## Passively supported

- [sc]GW[0]
- Damped
- All (Algo)

## Unsupported

- G-space projection
- NCORE > 1
- EFIELD\_PEAD





# Traditional DFT

## You should

- Run with MPS (multi-process service)
- Experiment with multiple CPU ranks per GPU

## Works best

- Large numbers of bands
- Large numbers of plane-waves

You can expect 2-4x for large systems with CPU/GPU balance; better on GPU-heavy workstations.



# Wait, what is MPS?

## Hyper-Q only works for kernels from the same context



- One context per MPI rank
- Without MPS, can't run multiple ranks concurrently
- This is a frustrating limitation

## MPS sits between MPI ranks and the GPUs

- Allows MPI ranks to 'share' a context
- Important when the kernels are small

# Example: Si super-cell

- 512 Si atoms
- 1282 bands
- 864000 PWs
- Algo = Normal





# Hybrid functionals (exact-exchange)

## You should

- Use 1 or 2 CPUs rank per GPU
- Set  $\text{NSIM} = \text{NBAND} / (2 * \text{NCPU})$



## Works best

- Large numbers of plane-waves
- Small number of ionic types

You can expect 1.5-6x, highly dependent on system size; better on GPU-heavy workstations.

# Example: $\beta$ -rhombohedral boron

- 105 Boron atoms
- 216 bands
- 110592 PWs
- Algo = Normal





# Road-map: Features

1. Gamma-point for very large unit cells
2. G-space projection for small to medium unit cells
3. Van der Waals density functional (vdF-DF)
4. Random phase approximation (RPA)
5. Active support for [sc]GW[0]
6. NCORE  $> 1$  for highly parallel runs



# Road-map: Performance

- Better performance for moderate sizes
  - Add blocking to all core kernels
  - Add batching to all library calls
- Better performance for large sizes
  - Update Magma support
  - Merge with threaded code base to reduce ranks per GPU
- Better performance for hybrid functionals
  - Parallelize outer loops
  - Pad projection sizes



# Road-map: Performance

- Better performance for moderate sizes
  - Add blocking to all core kernels
  - Add batching to all library calls
- Better performance for large sizes
  - Update Magma support
  - Merge with threaded code base to reduce ranks per GPU
- Better performance for hybrid functionals
  - Parallelize outer loops
  - Pad projection sizes





# Road-map: Performance

- Better performance for moderate sizes
  - Add blocking to all core kernels
  - Add batching to all library calls
- Better performance for large sizes
  - Update Magma support
  - Merge with threaded code base to reduce ranks per GPU
- Better performance for hybrid functionals
  - Parallelize outer loops
  - Pad projection sizes



# Road-map: Performance

- Better performance for moderate sizes
  - Add blocking to all core kernels
  - Add batching to all library calls
- Better performance for large sizes
  - Update Magma support
  - Merge with threaded code base to reduce ranks per GPU
- Better performance for hybrid functionals
  - Parallelize outer loops
  - Pad projection sizes



# Summary

GPU VASP will give you the right answer

- Extensive testing in Beta and for Vienna's acceptance

GPU VASP will give 2-4x performance on moderate to large systems

- The bigger the better

We are continuing to add feature support and improve performance

- Gamma-point is next on the list

Try it out.



# Summary

GPU VASP will give you the right answer

- Extensive testing in Beta and for Vienna's acceptance

GPU VASP will give 2-4x performance on moderate to large systems

- The bigger the better

We are continuing to add feature support and improve performance

- Gamma-point is next on the list

Try it out.



# Summary

GPU VASP will give you the right answer

- Extensive testing in Beta and for Vienna's acceptance

GPU VASP will give 2-4x performance on moderate to large systems

- The bigger the better

We are continuing to add feature support and improve performance

- Gamma-point is next on the list

Try it out.



# Summary

GPU VASP will give you the right answer

- Extensive testing in Beta and for Vienna's acceptance

GPU VASP will give 2-4x performance on moderate to large systems

- The bigger the better

We are continuing to add feature support and improve performance

- Gamma-point is next on the list

Try it out.



# Summary

GPU VASP will give you the right answer

- Extensive testing in Beta and for Vienna's acceptance

GPU VASP will give 2-4x performance on moderate to large systems

- The bigger the better

We are continuing to add feature support and improve performance

- Gamma-point is next on the list

**Try it out.**



# HOWTO: Compile

0. [Install CUDA toolkit]
1. Copy `arch/makefile.include.linux_intel_cuda` to `makefile.include`
2. Set most members as for the CPU build
3. Set `CUDA_ROOT` to the CUDA toolkit install location
4. Set `GENCODE_ARCH` to include GPU hardware capability
5. Set `MPI_INC` to point to MPI include files (`mpif90 --show`)
6. Build with `make gpu [gpu_ncl]`



# HOWTO: Run



1. Check NSIM
2. Normal mpirun with
  - 2-4 ranks per GPU for LDA/GGA
  - 1 rank per GPU for exact-exchange

# HOWTO: MPS



- Ideally, someone has set it up for you
  - On Titan, Blue Waters: `$ export CRAY_CUDA_MPS=1`
- Its a bit trickier otherwise

```
export CUDA\_VISIBLE\_DEVICES\=0
export CUDA_MPS_PIPE_DIRECTORY=/tmp/nvidia-mps
export CUDA_MPS_LOG_DIRECTORY=/tmp/nvidia-log
nvidia-cuda-mps-control -d
<run applications>
echo quit | nvidia-cuda-mps-control
```



# More performance





# What is VASP?

VASP is a complex package for performing ab-initio quantum-mechanical molecular dynamics (MD) simulations using pseudopotentials or the projector-augmented wave method and a plane wave basis set<sup>4</sup>.

---

<sup>4</sup>VASP the GUIDE

# Why VASP?

12-20% of CPU cycles @ HPC centers




## Academia

- Physics
- Physical chemistry
- Materials science
- Chemical engineering

## Industry

- Materials
- Big semiconductor
- Oil and gas
- Chemicals

## Usage @ Ohio SC's Oakley <sup>5</sup>



<sup>5</sup>12/14 – 2/15, via pbsacct